

Hobbeast — Business Request Summary — v1.16.0

Date: 2026-08-25 | Release: 1.16.0 | Theme: full source registry, dedup, restored admin

Business request

Register every source from the new V9 master list as a scraping destination (no duplicates), run the scraper on them with complete event data (description, image, venue, links, ticket link and price), prefer master/original sources over re-aggregating the same events from many mirrors, build event de-duplication, and restore the owner's admin surface with a view of destinations and daily/total scrape output.

What was delivered

- All 354 unique V9 sources are registered in the database registry with zero duplicates; the earlier 185 kept their state.
- The scraper is now registry-driven: sources rotate automatically (master ticket offices and national aggregators first), and each run logs its per-source results.
- Events are captured with rich data: description, image, venue/address, city, category tag for search, the ticket purchase link and the ticket price where published.
- Cross-source de-duplication is live: a fingerprint of the normalized title + date recognises the same real-world event seen through different sites, and the first (master) source wins. Proven live: 6 of 9 events from overlapping listings were correctly skipped as duplicates.
- The owner's admin access is restored with durable roles, and a new Admin > Scraper tab shows every destination, its last run and event counts, plus a daily breakdown and lifetime totals.

How it operates

GitHub Actions runs the scraper 3x daily for free; ~40 sources per run means the whole 354-source list is swept roughly every 3 days, with masters visited most often. Sources that yield nothing are visible in the admin dashboard and can be tuned or disabled without code changes.